Categorical variables

The Classification task introduced us to a type of non-numeric variable called Categorical.

A categorical variable

- Has a value drawn from a discrete set called *Categories* or *Classes*
    - The variable that is the target of a Classification task is Categorical
    - Hence the term "Classification" when the target is categorical

There is **no** ordering relationship between category/class values

- { "Red", "Green", "Blue" } (set notation)
- Versus *ordinal* values ["Small", "Medium", "Large"] (sequence notation)

There is no *magnitude* associated with a categorical value

- Even if we could order the colors: how much greater is "Blue" than "Red" ?

We will use C to denote the set of possible values in a category/class.

Since the values in C are unordered, C is mathematically a set of values
$$C = \{c_1, c_2, \ldots, \}$$

Since values in a category/class aren't ordered, they are often non-numeric.

Because computers deal with numbers, we will need to *encode* categorical variables as numbers.

In our Titanic example for Binary Classification, there were two obvious categorical variables

- `Survived` (the target)
- `Sex`

It might have gone unnoticed that the target was categorical

- Because the values were given to us encoded as numeric 1 (Survived) and 0 (not Survived)

We certainly did notice that Sex was non-numeric

- Because of it's encoding as text.

Our point is: don't count on the encoding of raw data in order to determine whether a variable is Categorical

We will illustrate this point with the `Pclass` variable, which has three possible distinct values.

How should we encode a Categorical variable with distinct values from a class C where $||C|| > 2$?

An obvious choice for such a variable is to encode the values with distinct integers.

This is usually a **bad** idea !

The `Pclass` feature was presented to us encoded as consecutive integers in $\{1, 2, 3\}$

But it could have just as easily been presented encoded as

- { "First", "Second", "Third" }.
- or $\{1, 2, 4\}$

Why is the encoding as $\{1, 2, 3\}$ any more correct than the encoding as $\{1, 2, 4\}$?

We will give a fuller answer in the module on Model Interpretation. For now:

- In a linear model
$$\hat{\mathbf{y}} = \Theta^T \mathbf{x}$$

  - Thus, the contribution of the j^{th} feature $\mathbf{x}_j$ to prediction $\hat{\mathbf{y}}$ is $\Theta_j * \mathbf{x}_j$

- Consider the encoding of $\mathbf{x}_j$ (`Pclass`) as $\{1, 2, 3\}$

  - The difference in contribution betwen "First", "Second" and "Third" are all equal
- Consider the encoding of $\mathbf{x}_j$ (`Pclass`) as $\{1, 2, 4\}$
  - The difference in contribution betwen "Second" and "Third" is twice that of "First" and "Second"

The arbitrary choice of encoding may have an impact on the prediction.

Bottom line

- Consider whether a feature should be treated as categorical *regardless* of the encoding presented
- Arbitrary mapping of a categorical value to an integer has consequences
    - Avoid it !

We will describe the proper way to encode categorical variables

- And revisit the Titanic example, changing `Pclass` from integer to categorical

One hot encoding (OHE)

So how should we encode a Categorical variable ?

We can encode a Categorical variable with values in class C with a vector $\mathbf{v}$ of length $||C||$.

- $\mathbf{v}$ will be all zero *except* at a single index j where $\mathbf{v}_j = 1$.

Consider

$$C = \{c_1, c_2, \ldots, c_{||C||}\}$$

The Categorical value c_k is encoded as vector $\mathbf{v}$

$$\mathbf{v}_j = 1 \quad \text{if } j = k$$
$$\mathbf{v}_j = 0 \quad \text{if } j \neq k$$

That is: position k of vector $\mathbf{v}$ is associated with the property: "Is value c_k"

Hence at most one position of the vector is non-zero.

Let's see what $\mathbf{v}$ looks like on a number of examples, one for each possible class $c \in C$:

	$\mathbf{v}_1$	$\mathbf{v}_2$	$\mathbf{v}_3$	$\cdots$	$\mathbf{v}_{||C||}$
c_1	1	0	0		0
c_2	0	1	0		0
c_3	0	0	1		0
$\cdots$					
$c_{||C||}$	0	0	0	$\cdots$	1

The vector representation essentially replaces the Categorical variable with $||C||$ binary variables $\mathbf{v}_1, \ldots, \mathbf{v}_{||C||}$

- Each an *indicator* or *dummy* variable: $\mathbf{v}_k$ indicates whether the value is c_k or not

This is called the **one hot encoding (OHE)** of a Categorical variable.

- Because at most one indicator in the representation is non-zero

We can use OHE on Categorical variables, whether they be targets or features.

To be concrete let's consider an example $\mathbf{x}$

- one numeric feature $\mathbf{x}_0$
- categorical feature $\mathbf{x}_1$ ("Gender") from categories $C_{(1)} = \{\mathrm{Male}, \mathrm{Female}\}$
- categorical feature $\mathbf{x}_2$ ("Location")from categories $C_{(2)} = \{\mathrm{Urban}, \mathrm{Suburban}, \mathrm{Exurban}\}$

And a few rows from our data set

$$\mathbf{X}' = \begin{pmatrix} \textbf{const} & \textbf{Gender} & \textbf{Location} \\ 1 & \text{Female} & \text{Urban} \\ 1 & \text{Female} & \text{Exurban} \\ 1 & \text{Male} & \text{Exurban} \\ 1 & \text{Male} & \text{Suburban} \\ & \vdots & \end{pmatrix}$$

After our One Hot Encoding we get

$$
\mathbf{X''}
=
\begin{pmatrix}
\textbf{const} & \textbf{IsFemale} & \textbf{IsMale} & \textbf{IsUrban} & \textbf{IsSuburban} & \textbf{IsExurban} \\
1 & 1 & 0 & 1 & 0 & 0 \\
1 & 1 & 0 & 0 & 0 & 1 \\
1 & 0 & 1 & 0 & 0 & 1 \\
1 & 0 & 1 & 0 & 1 & 0 \\
\vdots & & & & &
\end{pmatrix}
$$

Notice that the number of features has increased.

Specifically: a single categorical variable x_j from a category $C = \{c_1, c_2, \ldots\}$ of size $||C||$

- has been replaced by $||C||$ new binary, indicator variables
    - Is c_1
    - Is c_2
    - $\vdots$
    - Is $c_{||C||}$

Categorical features versus categorical targets

Although OHE can be applied to features $\mathbf{x}$ or targets $\mathbf{y}$, there are some subtle differences in practice

Categorical targets

Although we should use OHE to encode the targets, *in practice* you might see targets encoded as integers

- Binary targets as 0/1
- Other targets as integers
  - `sklearn` method `LabelEncoder` does exactly this

If it's such a bad idea: why does this happen ?

The answer

- It **may** not matter *from a coding perspective*
    - Often, the code need only be able to *distinguish between* target values
        - e.g., restrict the examples to those with a particular value of the target
    - So the encoding of values is not important
    - In fact: `sklearn` is perfectly happy with non-numeric targets for just this reason !

It **may matter** from a mathematical perspective

- Negative/Positive will often be encoded by either $0/1$ or $-1/+1$
- For example: when Negative/Positive encoding is $-1/+1$

$$10 + \mathbf{y}^{(i)} \;=\; 9 \quad \text{for Negative example } i$$

$$10 + \mathbf{y}^{(i)} \;=\; 11 \quad \text{for Positive example } i$$

You will often see Categorical values manipulated as mathematical objects when they are used to define Loss Functions.

So please be aware of the purpose for which you are encoding.

Categorical features

We would love to make the blanket statement: Always use OHE for categorical features.

Unfortunately, there is one model in which OHE may cause a problem

- Linear Regression, with an intercept
- There is a simple fix (i.e., an argument to pass to implementations of OHE)

The issue is called the *Dummy variable* trap and will be explained in a subsequent module.

Text: another categorial variables

How do you include text variables ? One-hot encoding of the vocabulary !

Example: Spam filtering (Classification task with target: Is Spam/Is Not Spam)

In theory, OHE is the solution

- Vocabulary V of possible words
- $||V||$ indicator variables

In practice

- Vocabulary can be big ! Lots of variables, lots of memory required using OHE
- The representation of a word is "sparse": a single 1 and lots of 0's
    - no relationship between related words: dog, dogs
- Lots of feature engineering possibilities: an ALL CAP feature

We will devote a subsequent module to the topic of Natural Language Processing.

Recap

- We introduced methods to deal with non-numeric variables
- Unfortunately, there are some nuances

```python
In [4]: print("Done")
```

Done